


MongoDB III

CS 341 Database Systems

Basic Comparison



Schema Design Best Practices

- Usually, flexible
- Major Considerations
 - Storage of data
 - Good query performance
 - Reasonable amount of hardware

Relations

- **One-to-One** - Prefer key value pairs within the document
 - Department and Manager
- **One-to-Few** - Prefer embedding
 - Customer and Addresses
- **One-to-Many** - Prefer embedding
 - Department and Staff
- **Many-to-Many** - Referencing or 1-way embedding or 2-way embedding
 - Courses and Students

Consider a Scenario

- Each SalesRep is assigned to one department. Each department has many sales representatives
- Each department is headed by a manager who is also a SalesRep
- Each customer can be dealt by multiple SalesRep and different SalesRep can deal with different customers

Consider a Scenario

- Each **SalesRep** is assigned to one **department**. Each department has many sales representatives
- Each department is headed by a manager who is also a SalesRep
- Each **customer** can be dealt by multiple SalesRep and different SalesRep can deal with different customers

Department - SalesRep
1-M (Assigned)

Department - SalesRep
1-1 (Manager)

Customer - SalesRep
M-M (Deals)

Collections

- **Department**
 - SalesRep as key:value pair or embedded document (Manager)
- **SalesRep**
 - Department as embedded document
 - Embedded array of customers dealt by this salesRep
- **Customer**
 - Embedded array of SalesRep dealing this customer

Refer to the scenario create collections with appropriate validation rules.

- Each SalesRep is assigned to one department. Each department has many sales representatives.
- Each department is headed by a manager who is also a SalesRep
- Each customer can be dealt by multiple SalesRep and different SalesRep can deal with different customers

The Sales Rep collection stores the `s_name`, `s_email` and `s_performance` score. The customer collection contains the `cust_name`, `cust_add` and `cust_email`. The department collection contains the `dept_name`, `dept_region`, and grade.

Limitations:

- A SalesRep can only deal with maximum 10 customers
- A customer is dealt by maximum 5 SalesRep.
- The `performance_score` is given as a number between 1 and 10.
- The `dept_region` must be from among East, West, Central
- Grade is a single character string.
- The fields in red are mandatory fields.
- To implement the M-M relation, both sides will keep embedded documents. We only require name and email address in the array. *Hint: sub-schema*

Departments

```
db.createCollection("departments", {  
    validator: {  
        $jsonSchema: {  
            bsonType: "object",  
            required: ["dept_name", "dept_region", "dept_manager"],  
            properties: {  
                dept_name: {  
                    bsonType: "string",  
                    description: "Department name is required and must be unique."  
                },  
                dept_region: {  
                    bsonType: "string",  
                    enum: ["East", "West", "Central"],  
                    description: "Department region must be East, West, or Central."  
                },  
                grade: {  
                    bsonType: "string",  
                    maxLength: 1,  
                    description: "Grade must be a single character."  
                },  
            }  
        }  
    }  
});
```

```
        dept_manager: {  
            bsonType: "object",  
            required: ["s_name", "s_email"],  
            properties: {  
                s_name: {  
                    bsonType: "string",  
                    description: "Manager's name is required."  
                },  
                s_email: {  
                    bsonType: "string",  
                    description: "Manager's email is required."  
                },  
                description: "Manager who is also a SalesRep."  
            }  
        }  
    }  
});
```

Scenario



Refer to the scenario create collections with appropriate validation rules.

- Each SalesRep is assigned to one department. Each department has many sales representatives.
- Each department is headed by a manager who is also a SalesRep
- Each customer can be dealt by multiple SalesRep and different SalesRep can deal with different customers

The Sales Rep collection stores the `s_name`, `s_email` and `s_performance` score. The customer collection contains the `cust_name`, `cust_add` and `cust_email`. The department collection contains the `dept_name`, `dept_region`, and grade.

Limitations:

- A SalesRep can only deal with maximum 10 customers
- A customer is dealt by maximum 5 SalesRep.
- The `performance_score` is given as a number between 1 and 10.
- The `dept_region` must be from among East, West, Central
- Grade is a single character string.
- The fields in red are mandatory fields.
- To implement the M-M relation, both sides will keep embedded documents. We only require name and email address in the array. *Hint: sub-schema*

```
db.createCollection("salesreps", {  
    validator: {  
        $jsonSchema: {  
            bsonType: "object",  
            required: ["s_name", "s_email", "departmentId"],  
            properties: {  
                s_name: {  
                    bsonType: "string",  
                    description: "SalesRep name is required."  
                },  
                s_email: {  
                    bsonType: "string",  
                    description: "SalesRep email is required unique."  
                },  
                s_performance: {  
                    bsonType: "number",  
                    minimum: 1,  
                    maximum: 10,  
                    description: "SalesRep performance score must be a  
number between 1 and 10."  
                },  
                departmentId: {  
                    bsonType: "objectId",  
                    description: "SalesRep must be assigned to a department."  
                },  
            }  
        }  
    }  
});
```

Linked

SalesRep

```
customers: {  
    bsonType: "array",  
    maxItems: 10,  
    items: {  
        bsonType: "object",  
        required: ["cust_name", "cust_email"]  
        properties: {  
            cust_name: {  
                bsonType: "string",  
                description: "Customer name is required."  
            },  
            cust_email: {  
                bsonType: "string",  
                description: "Customer email is required."  
            }  
        }  
    },  
    description: "List of customers associated with  
the sales rep."  
}  
};
```

Array of embedded documents.



SalesRep

```
db.createCollection("salesreps", {  
    validator: {  
        $jsonSchema: {  
            bsonType: "object",  
            required: ["s_name", "s_email", "department"],  
            properties: {  
                s_name: {  
                    bsonType: "string",  
                    description: "SalesRep name is required."  
                },  
                s_email: {  
                    bsonType: "string",  
                    description: "SalesRep email is required."  
                },  
                s_performance: {  
                    bsonType: "number",  
                    minimum: 1,  
                    maximum: 10,  
                    description: "SalesRep performance score must be a  
number between 1 and 10."  
                },  
                department: {  
                    bsonType: "string",  
                    description: "SalesRep must be assigned to a department."  
                },  
            }  
        }  
    }  
}
```

Key-value

Database Systems - Abeera Tariq

SalesRep

Embedded

Scenario



Refer to the scenario create collections with appropriate validation rules.

- Each SalesRep is assigned to one department. Each department has many sales representatives.
- Each department is headed by a manager who is also a SalesRep
- Each customer can be dealt by multiple SalesRep and different SalesRep can deal with different customers

The Sales Rep collection stores the `s_name`, `s_email` and `s_performance` score. The customer collection contains the `cust_name`, `cust_add` and `cust_email`. The department collection contains the `dept_name`, `dept_region`, and grade.

Limitations:

- A SalesRep can only deal with maximum 10 customers
- A customer is dealt by maximum 5 SalesRep.
- The `performance_score` is given as a number between 1 and 10.
- The `dept_region` must be from among East, West, Central
- Grade is a single character string.
- The fields in red are mandatory fields.
- To implement the M-M relation, both sides will keep embedded documents. We only require name and email address in the array. *Hint: sub-schema*

Customers

```
db.createCollection("customers", {  
  validator: {  
    $jsonSchema: {  
      bsonType: "object",  
      required: ["cust_email"],  
      properties: {  
        cust_name: {  
          bsonType: "string",  
          description: "Customer name is required."  
        },  
        cust_addr: {  
          bsonType: "string"  
        },  
        cust_email: {  
          bsonType: "string",  
          description: "Customer email is required."  
        },  
        salesReps: {  
          bsonType: "array",  
          maxItems: 5,  
          items: {  
            bsonType: "object",  
            required: ["s_name", "s_email"],  
            properties: {  
              s_name: {  
                bsonType: "string",  
                description: "SalesRep name is required."  
              },  
              s_email: {  
                bsonType: "string",  
                description: "SalesRep email is required."  
              }  
            }  
          }  
        }  
      }  
    }  
  }  
});
```

```
  s_email: {  
    bsonType: "string",  
    description: "SalesRep email is required."  
  }  
}  
}  
},  
description: "List of sales reps associated with the  
customer."  
}  
}  
}  
}  
});
```

Querying Linked Documents



SalesRep

```
db.createCollection("salesreps", {  
  validator: {  
    $jsonSchema: {  
      bsonType: "object",  
      required: ["s_name", "s_email", "departmentId"],  
      properties: {  
        s_name: {  
          bsonType: "string",  
          description: "SalesRep name is required."  
        },  
        s_email: {  
          bsonType: "string",  
          description: "SalesRep email is required unique."  
        },  
        s_performance: {  
          bsonType: "number",  
          minimum: 1,  
          maximum: 10,  
          description: "SalesRep performance score must be a  
number between 1 and 10."  
        },  
        departmentId: {  
          bsonType: "objectId",  
          description: "SalesRep must be assigned to a department."  
        },  
      }  
    }  
  }  
});
```

```
customers: {  
  bsonType: "array",  
  maxItems: 10,  
  items: {  
    bsonType: "object",  
    properties: {  
      cust_name: {  
        bsonType: "string",  
        description: "Customer name is required."  
      },  
      cust_email: {  
        bsonType: "string",  
        description: "Customer email is required."  
      }  
    }  
  },  
  description: "List of customers associated with  
the sales rep."  
}  
};
```

\$lookup

- The **\$lookup** operator allows you to perform a **join** between two collections in MongoDB.

```
db.salesreps.aggregate([
  {
    $lookup: {
      from: "departments", //The collection to join with
      localField: "departmentId", //The field in `salesreps` that stores the department's ObjectId
      foreignField: "_id", //The field in `departments` that contains the ObjectId
      as: "departmentInfo" //The field where the resulting department information will be stored
    }
  },
  {
    $unwind: "$departmentInfo" // Flatten the array created by `$lookup` so each sales rep gets a
                                single department
  }
]);
```

SALES REP

```
{  
  "_id": ObjectId("..."),  
  "s_name": "John Doe",  
  "s_email": "john.doe@example.com",  
  "departmentId":  
    ObjectId("60b8a9bb1d52c0d2f7a8c8d9")  
}
```

DEPARTMENT

```
{  
  "_id": ObjectId("60b8a9bb1d52c0d2f7a8c8d9"),  
  "dept_name": "Sales",  
  "dept_region": "East",  
  "grade": "A",  
  "manager_name": "Alice Smith"  
}
```

After \$lookup stage

```
{  
  "_id": ObjectId("..."),  
  "s_name": "John Doe",  
  "s_email": "john.doe@example.com",  
  "departmentId":  
  ObjectId("60b8a9bb1d52c0d2f7a8c8d9"),  
  "departmentInfo": [  
    {  
      "_id": ObjectId("60b8a9bb1d52c0d2f7a8c8d9"),  
      "dept_name": "Sales",  
      "dept_region": "East",  
      "grade": "A",  
      "manager_name": "Alice Smith"  
    }  
  ]  
}
```

After \$unwind stage

```
{  
  "_id": ObjectId("..."),  
  "s_name": "John Doe",  
  "s_email": "john.doe@example.com",  
  "departmentId":  
  ObjectId("60b8a9bb1d52c0d2f7a8c8d9"),  
  "departmentInfo": {  
    "_id":  
    ObjectId("60b8a9bb1d52c0d2f7a8c8d9"),  
    "dept_name": "Sales",  
    "dept_region": "East",  
    "grade": "A",  
    "manager_name": "Alice Smith"  
  }  
}
```

Orders and Inventory

```
db.orders.insertMany([  
  { "_id" : 1, "item" : "almonds", "price" : 12,  
    "quantity" : 2 },  
  { "_id" : 2, "item" : "pecans", "price" : 20, "quantity" :  
    1 },  
  { "_id" : 3 }  
])
```


```
db.inventory.insertMany([  
  { "_id" : 1, "sku" : "almonds", "description": "product 1",  
    "instock" : 120 },  
  { "_id" : 2, "sku" : "bread", "description": "product 2",  
    "instock" : 80 },  
  { "_id" : 3, "sku" : "cashews", "description": "product 3",  
    "instock" : 60 },  
  { "_id" : 4, "sku" : "pecans", "description": "product 4",  
    "instock" : 70 },  
  { "_id" : 5, "sku": null, "description": "Incomplete" },  
  { "_id" : 6 }  
])
```

```
db.orders.aggregate( [  
  {  
    $lookup:  
      {  
        from: "inventory",  
        localField: "item",  
        foreignField: "sku",  
        as: "inventory_docs"  
      }  
  }  
)
```

Output

```
{  
  "_id" : 1,  
  "item" : "almonds",  
  "price" : 12,  
  "quantity" : 2,  
  "inventory_docs" : [  
    { "_id" : 1, "sku" : "almonds", "description" : "product 1",  
    "instock" : 120 }  
  ]  
}  
  
{  
  "_id" : 2,  
  "item" : "pecans",  
  "price" : 20,  
  "quantity" : 1,  
  "inventory_docs" : [  
    { "_id" : 4, "sku" : "pecans", "description" : "product 4", "instock" : 70 }  
  ]  
}  
  
{  
  "_id" : 3,  
  "inventory_docs" : [  
    { "_id" : 5, "sku" : null, "description" : "Incomplete" },  
    { "_id" : 6 }  
  ]  
}
```

Consider a Relational Schema translated to Collections → *The Linking Approach*



student (sid, sname, byear)
course (cid, cname, teacher)
take (sid, cid)

student (sid, sname, byear)
course (cid, cname, teacher)
take (sid, cid)

Find the course IDs of all courses taken by the student with the name "Raymond"

It involves a "join" operation between two collections.

```
db.take.aggregate(  
[  
  {  
    $lookup: {  
      localField: "sid",  
      from: "student",  
      foreignField: "sid",  
      as: "student_info"  
    }  
  },  
  { $unwind: "$student_info"},  
  { $match: { "student_info.sname": "Raymond" } },  
  { $project: { cid: 1, _id: 0 } }  
]
```

Natural Join

```
SELECT cid  
FROM take JOIN student ON take.sid=student.sid  
WHERE student.sname = "Raymond"
```

Output

```
{ cid : "90123" }  
{ cid : "90696" }  
{ cid : "95270" }
```



In the output, there can be multiple occurrences of the same course ID.

It could be regarded as “correct” since we do not require “distinct” course IDs in the output.

Suppose that we require the “distinct” requirement. What should we do?

```
db.take.aggregate(  
[  
  {  
    $lookup: {  
      localField: "sid",  
      from: "student",  
      foreignField: "sid",  
      as: "student_info"  
    }  
  },  
  { $unwind: "$student_info"},  
  { $match: { "student_info.sname": "Raymond" } },  
  { $project: { cid: 1, _id: 0 } },  
  { $group: { _id: "$cid" } }  
])
```

Find the student IDs of all students who take the course taught by “Prof. Charles”

student (sid, sname, byear)
course (cid, cname, teacher)
take (sid, cid)



```
db.take.aggregate(  
[  
  {  
    $lookup: {  
      localField: "cid",  
      from: "course",  
      foreignField: "cid",  
      as: "course_info"  
    }  
  },  
  { $unwind: "$course_info"},  
  { $match: { "course_info.teacher": "Prof. Charles" } },  
  { $project: { sid: 1, _id: 0 } }  
]  
)
```

Names of all students taking the course by “Prof. Charles”

```
db.student.aggregate([
  {
    $lookup: {
      localField: "sid",
      from: "take",
      foreignField: "sid",
      as: "student_take"
    }
  },
  { $unwind: "$student_take" },
  {
    $lookup: {
      localField: "student_take.cid",
      from: "course",
      foreignField: "cid",
      as: "student_take_course",
    }
  },
  { $unwind: "$student_take_course" },
  { $match: { "student_take_course.teacher": "Prof. Charles" } },
  { $project: { sname: 1, _id: 0} }
])
```

For distinct values, add another stage
{ \$group: { _id: "\$sname" }}

**Find the class size of each course.
Show the course ID, the course name
and the class size for each course.**

```
db.take.aggregate(  
[  
  { $group : { _id : "$cid", classSize : { $sum : 1 } } },  
  { $lookup: {  
    localField: "_id", // from previous stage  
    from: "course",  
    foreignField: "cid",  
    as: "course_info"  
  } },  
  { $unwind: "$course_info" },  
  { $project: { _id: 1, cname: "$course_info.cname", classSize: 1 } } // Rename course name  
]  
);
```

student (sid, sname, byear)
course (cid, cname, teacher)
take (sid, cid)



Computing Values

\$add
\$subtract
\$multiply
\$divide

Output the field "sid" and a new field called "newSum" which is equal to "a+2" for each document

```
{sid: 1, a:100, b:150}  
{sid: 2, a:70, b:200}  
{sid: 3, a:500, b:300}
```



```
db.temp4.aggregate([  
  { $project: {newSum: {$add: ["$a", 2]}, sid: 1, _id: 0 } }  
])
```

Output

```
{ sid : 1, newSum : 102 }  
{ sid : 2, newSum : 72 }  
{ sid : 3, newSum : 502 }
```

Output the field “sid” and a new field called “newProduct” which is equal to 2 times “a” for each document

```
{sid: 1, a:100, b:150}  
{sid: 2, a:70, b:200}  
{sid: 3, a:500, b:300}
```

```
db.temp4.aggregate([  
  { $project: {newProduct: {$multiply: ["$a", 2]}, sid: 1, _id: 0 } }  
])
```

Output

```
{ sid : 1, newProduct : 200 }  
{ sid : 2, newProduct : 140 }  
{ sid : 3, newProduct : 1000 }
```