


# Transaction Management and Concurrency

CS 341 Database Systems



## ACCESS TO THE DATABASE

- Large databases are usually shared by many users, and resources
- It is efficient to allow concurrent access.

# Transaction

- “An action, or series of actions, carried out by a single user or application program, which **reads** or **updates** the contents of the **database**.”

# Transaction

- A transaction is a **logical unit of work**, as well as **unit of recovery**.
- It is broken down into a sequence of atomic operations, which if any fail, the whole transaction is undone.

SELECT | INSERT | ...

...

... work ...

...

COMMIT | ROLLBACK

# Transactions

- Series of database commands with clear semantics
  - e.g. transfer of funds from one account to another
- **Commit:**
  - If nothing fails, **commit** point where the DB should be consistent.
  - All updates are **tentative** until committed.
- **Rollback:**
  - If any command fails → whole series of commands is undone.
  - Any DBMS support these (and language e.g. SQL)

# Transaction


A transaction is any action that **reads** from and/or **writes** to a database.

A transaction may consist of the following:

- A simple **SELECT** statement to generate a list of table contents.
- A series of related **UPDATE** statements to change the values of attributes in various tables.
- A series of **INSERT** statements to add rows to one or more tables.
- A combination of **SELECT**, **UPDATE**, and **INSERT** statements.

- Suppose that you sell a product to a customer and the customer may charge the purchase to his or her account.
- The sales transaction consists of at least the following parts:
  - Write a new customer invoice.
  - Reduce the quantity on hand in the product's inventory.
  - Update the account transactions.
  - Update the customer balance.

**Figure 10.1** The Ch10\_SaleCo Database Relational Diagram



INSERT INTO INVOICE

VALUES (1009, 10016, '18-Jan-2022', 256.99, 20.56, 277.55, 'cred', 0.00, 277.55);

INSERT INTO LINE

VALUES (1009, 1, '89-WRE-Q', 1, 256.99, 256.99);

UPDATE PRODUCT

SET PROD\_QOH = PROD\_QOH - 1

WHERE PROD\_CODE = '89-WRE-Q';

UPDATE CUSTOMER

SET CUST\_BALANCE = CUST\_BALANCE + 277.55


WHERE CUST\_NUMBER = 10016;

INSERT INTO ACCT\_TRANSACTION

VALUES (10007, '18-Jan-22', 10016, 'charge', 277.55);

COMMIT;


# ACID Properties



# ACID Properties

- **Atomicity:** all or nothing (any error → Rollback, as if nothing happened)
- **Consistency:** a consistent state always leads to another consistent state
- **Isolation:** a transaction's updates are hidden until it Commits
- **Durability:** after a Commit, updates persist
- These are the **ACID** properties of transactions.

# Transaction State



# Transaction State

- **Active** - the initial state; the transaction stays in this state while it is executing
- **Partially committed** - after the final statement has been executed.
- **Failed** - after the discovery that normal execution can no longer proceed.
- **Aborted** - after the transaction has been rolled back and the database restored to its state prior to the start of the transaction.  
Two options after it has been aborted:
  - Restart the transaction - Can be done only if no internal logical error
  - Kill the transaction
- **Committed** - after successful completion.

# System Recovery


- How does the system recover after a system failure (e.g., power failure) or media failure (e.g., disk crash)?
- In the event of a crash ...
  - Contents of main memory are lost.
  - Transaction log persists.
  - At failure, certain transactions will be complete while others partially complete.
- Note that updates are held in memory buffers and written out periodically.

# Recovery

To recover the state of the database we can use:

- A **log file** recording every database operation.
- **Checkpoints** recording the state of all active transactions.
  - Then: develop an algorithm for transactions to **UNDO**,
  - and those that we need to **REDO**, to effect recovery.
  - at intervals, the system will:
    - Flush its buffers – buffers are forced to write changes to secondary storage.
    - Write out a checkpoint record to log indicating which transactions are in progress.

# Five Transaction categories



- The most recent check point record was taken at time  $tc$ .

# Transactions

Completed	Un-Finished
Cached	Written
T2 T4	T1
	T3 T5

# Transactions

- CHECKPOINT RECORD:
  - T3, T5 : **undone** (rollback possible)
  - T2, T4 : **re-done**
- DBMS creates REDO/UNDO list from checkpoint record + system log.
- **Isolation**
  - order of recovery *not* crucial,
  - only DB should be consistent near tf (time of failure)

# Concurrency

# Concurrency

- Many users hence many transactions - at the same time.
- Databases are **shared!**
- So: Transactions must be **isolated** => need of **concurrency control** to ensure no interference.
- We will look at:
  - **3 classic problems** on concurrent access
  - **Locking** mechanism
  - **Deadlock** resolution

# Three classic problems

**Problem:** *Two or more transactions read / write on the same part of the database.*

Although transactions execute correctly, results may **interleave** in different ways leading to the 3 classic problems.

- *Lost Update*
- *Uncommitted Dependency*
- *Inconsistent Analysis*

# Lost Update problem

Time	User 1 (Trans A)	User2 (Trans B)
1	Retrieve t	
2		Retrieve t
3	Update t	
4		Update t
5		
6		
7		

t : tuple in a table.  
Trans A loses an update at t4.  
The update at t3 is lost (overwritten) at t4 by B.

# Lost Update problem

Time	T <sub>1</sub>	T <sub>2</sub>	bal <sub>x</sub>
t <sub>1</sub>		begin_transaction	100
t <sub>2</sub>	begin_transaction	read(bal <sub>x</sub> )	100
t <sub>3</sub>	read(bal <sub>x</sub> )	200    bal <sub>x</sub> = bal <sub>x</sub> + 100	100
t <sub>4</sub>	90    bal <sub>x</sub> = bal <sub>x</sub> - 10	200    write(bal <sub>x</sub> )	200
t <sub>5</sub>	90    write(bal <sub>x</sub> )	commit	90
t <sub>6</sub>	commit		90

- An apparently successfully completed update operation by one user can be overridden by another user. This is known as the lost update problem

# Serial VS Concurrent

**Table 10.3** Serial Execution of Two Transactions

Time	Transaction	Step	Stored Value
1	T1	Read PROD_QOH	35
2	T1	$\text{PROD\_QOH} = 35 + 100$	
3	T1	Write PROD_QOH	135
4	T2	Read PROD_QOH	135
5	T2	$\text{PROD\_QOH} = 135 - 30$	
6	T2	Write PROD_QOH	105

**Table 10.4** Lost Updates

Time	Transaction	Step	Stored Value
1	T1	Read PROD_QOH	35
2	T2	Read PROD_QOH	35
3	T1	$\text{PROD\_QOH} = 35 + 100$	
4	T2	$\text{PROD\_QOH} = 35 - 30$	
5	T1	Write PROD_QOH (lost update)	135
6	T2	Write PROD_QOH	5

# Uncommitted Dependency

- 2 Problems (T1-3 ; T6-8).
- One trans is allowed to retrieve/update) a tuple updated by another, but not yet committed.
- Trans A is dependent at time t2 on an **uncommitted change** made by Trans B, which is lost on Rollback.

Time	User 1 (Trans A)	User 2 (Trans B)
1		Update t
2	Retrieve t	
3		Rollback
4		
5		
6		Update t
7	Update t	
8		Rollback

# Uncommitted Dependency

Time	T <sub>3</sub>	T <sub>4</sub>	bal <sub>x</sub>
t <sub>1</sub>		begin_transaction	100
t <sub>2</sub>		read(bal <sub>x</sub> )	100
t <sub>3</sub>		bal <sub>x</sub> = bal <sub>x</sub> + 100	100
t <sub>4</sub>	begin_transaction	write(bal <sub>x</sub> )	200
t <sub>5</sub>	read(bal <sub>x</sub> )	:	200
t <sub>6</sub>	bal <sub>x</sub> = bal <sub>x</sub> - 10	rollback	100
t <sub>7</sub>	write(bal <sub>x</sub> )		190
t <sub>8</sub>	commit		190

- The uncommitted dependency problem occurs when one transaction is allowed to see the intermediate results of another transaction before it has committed.

**Table 10.5** Transactions Creating an Uncommitted Data Problem

Transaction	Computation
T1: Purchase 100 units	$\text{PROD\_QOH} = \text{PROD\_QOH} + 100$ (Rolled back)
T2: Sell 30 units	$\text{PROD\_QOH} = \text{PROD\_QOH} - 30$

**Table 10.6** Correct Execution of Two Transactions

Time	Transaction	Step	Stored Value
1	T1	Read PROD_QOH	35
2	T1	$\text{PROD\_QOH} = 35 + 100$	
3	T1	Write PROD_QOH	135
4	T1	*****ROLLBACK*****	35
5	T2	Read PROD_QOH	35
6	T2	$\text{PROD\_QOH} = 35 - 30$	
7	T2	Write PROD_QOH	5

**Table 10.7** An Uncommitted Data Problem

Time	Transaction	Step	Stored Value
1	T1	Read PROD_QOH	35
2	T1	$\text{PROD\_QOH} = 35 + 100$	
3	T1	Write PROD_QOH	135
4	T2	Read PROD_QOH (Read uncommitted data)	135
5	T2	$\text{PROD\_QOH} = 135 - 30$	
6	T1	***** ROLLBACK *****	35
7	T2	Write PROD_QOH	105

# Inconsistent Analysis

- Trans A sees inconsistent DB state after B updated Accumulator => performs inconsistent analysis.

Initially: Acc 1 = 40;      Acc2 = 50;      Acc3 = 30;

Time	User 1 (Trans A)	User 2 (Trans B)
1	Retrieve Acc 1 : Sum = 40	
2	Retrieve Acc2 : Sum = 90	
3		Retrieve Acc3 :
4		Update Acc3: 30 → 20
5		Retrieve Acc1:
6		Update Acc1: 40 → 50
7		commit
8	Retrieve Acc3: Sum = 110 (not 120)	

# Inconsistent Analysis

Time	T <sub>5</sub>	T <sub>6</sub>	bal <sub>x</sub>	bal <sub>y</sub>	bal <sub>z</sub>	sum
t <sub>1</sub>		begin_transaction	100	50	25	
t <sub>2</sub>	begin_transaction	sum = 0	100	50	25	0
t <sub>3</sub>	read(bal <sub>x</sub> )	read(bal <sub>x</sub> )	100	50	25	0
t <sub>4</sub>	→ bal <sub>x</sub> = bal <sub>x</sub> - 10	sum = sum + bal <sub>x</sub>	100	50	25	100
t <sub>5</sub>	→ write(bal <sub>x</sub> )	read(bal <sub>y</sub> )	90	50	25	100
t <sub>6</sub>	read(bal <sub>z</sub> )	sum = sum + bal <sub>y</sub>	90	50	25	150
t <sub>7</sub>	bal <sub>z</sub> = bal <sub>z</sub> + 10		90	50	25	150
t <sub>8</sub>	write(bal <sub>z</sub> )		90	50	35	150
t <sub>9</sub>	commit	read(bal <sub>z</sub> )	90	50	35	150
t <sub>10</sub>		sum = sum + bal <sub>z</sub>	90	50	35	185
t <sub>11</sub>		commit	90	50	35	185

The problem of inconsistent analysis occurs when a transaction reads several values from the database but a second transaction updates some of them during the execution of the first.

# Inconsistent Analysis Eg.2

**Table 10.8** Retrieval during Update

Transaction T1	Transaction T2
SELECT SUM(PROD_QOH) FROM PRODUCT	UPDATE PRODUCT SET PROD_QOH = PROD_QOH + 10 WHERE PROD_CODE = 1546-QQ2
	UPDATE PRODUCT SET PROD_QOH = PROD_QOH - 10 WHERE PROD_CODE = 1558-QW1
	COMMIT;

**Table 10.9** Transaction Results: Data Entry Correction

	Before	After
PROD_CODE	PROD_QOH	PROD_QOH
11QER/31	8	8
13-Q2/P2	32	32
1546-QQ2	15	(15 + 10) → 25
1558-QW1	23	(23 - 10) → 13
2232-QTY	8	8
2232-QWE	6	6
<b>Total</b>	<b>92</b>	<b>92</b>

# Inconsistent Analysis Eg.2

**Table 10.10 Inconsistent Retrievals**

Time	Transaction	Action	Value	Total
1	T1	Read PROD_QOH for PROD_CODE = '11QER/31'	8	8
2	T1	Read PROD_QOH for PROD_CODE = '13-Q2/P2'	32	40
3	T2	Read PROD_QOH for PROD_CODE = '1546-QQ2'	15	
4	T2	PROD_QOH = 15 + 10		
5	T2	Write PROD_QOH for PROD_CODE = '1546-QQ2'	25	
6	T1	Read PROD_QOH for PROD_CODE = '1546-QQ2'	25	(After) 65
7	T1	Read PROD_QOH for PROD_CODE = '1558-QW1'	23	(Before) 88
8	T2	Read PROD_QOH for PROD_CODE = '1558-QW1'	23	
9	T2	PROD_QOH = 23 - 10		
10	T2	Write PROD_QOH for PROD_CODE = '1558-QW1'	13	
11	T2	***** COMMIT *****		
12	T1	Read PROD_QOH for PROD_CODE = '2232-QTY'	8	96
13	T1	Read PROD_QOH for PROD_CODE = '2232-QWE'	6	102

# Read Write Conflicts

**Table 10.11** Read/Write Conflict Scenarios: Conflicting Database Operations Matrix

		Transactions		
		T1	T2	Result
Operations	Read	Read	No conflict	
	Read	Write	Conflict	
	Write	Read	Conflict	
	Write	Write	Conflict	

# Why these problems occur?

- Retrieve : 'read' (R)
- Update : 'write' (W).
- Interleaving two transactions:

**RR**      No problem

**WW**      Lost update

**WR**      Uncommitted dependency

**RW**      Inconsistent analysis

Time	T <sub>1</sub>	T <sub>2</sub>	bal <sub>x</sub>
t <sub>1</sub>			100
t <sub>2</sub>	begin_transaction		100
t <sub>3</sub>	read(bal <sub>x</sub> )		100
t <sub>4</sub>	bal <sub>x</sub> = bal <sub>x</sub> - 10		200
t <sub>5</sub>	write(bal <sub>x</sub> )		90
t <sub>6</sub>	commit	begin_transaction	90
		read(bal <sub>x</sub> )	100
		bal <sub>x</sub> = bal <sub>x</sub> + 100	200
		write(bal <sub>x</sub> )	200
		commit	190

Lost Update WW

Uncommitted Dependency WR

Time	T <sub>5</sub>	T <sub>6</sub>	bal <sub>x</sub>	bal <sub>y</sub>	bal <sub>z</sub>	sum
t <sub>1</sub>		begin_transaction	100	50	25	
t <sub>2</sub>	begin_transaction	sum = 0	100	50	25	0
t <sub>3</sub>	read(bal <sub>x</sub> )	read(bal <sub>x</sub> )	100	50	25	0
t <sub>4</sub>	bal <sub>x</sub> = bal <sub>x</sub> - 10	sum = sum + bal <sub>x</sub>	100	50	25	100
t <sub>5</sub>	write(bal <sub>x</sub> )	read(bal <sub>y</sub> )	90	50	25	100
t <sub>6</sub>	read(bal <sub>z</sub> )	sum = sum + bal <sub>y</sub>	90	50	25	150
t <sub>7</sub>	bal <sub>z</sub> = bal <sub>z</sub> + 10		90	50	35	150
t <sub>8</sub>	write(bal <sub>z</sub> )		90	50	35	150
t <sub>9</sub>	commit	read(bal <sub>z</sub> )	90	50	35	150
t <sub>10</sub>		sum = sum + bal <sub>z</sub>	90	50	35	185
t <sub>11</sub>		commit	90	50	35	185

Inconsistent Analysis RW

# Session 02

# ANSI/ISO Transaction Isolation Levels

- Transaction isolation levels refer to the degree to which transaction data is "**protected or isolated**" from other **concurrent** transactions.
- The isolation levels are described based on what data other transactions can see (read) during execution. More precisely, the *transaction isolation levels are described by the type of "reads" that a transaction allows or does not allow.*

# Types of Read Operations

## 1. Dirty reads

- A transaction reads data that has been written by another transaction that has not been committed yet.

## 2. Nonrepeatable (fuzzy) reads

- A transaction re-reads data it has previously read and finds that another committed transaction has *modified or deleted* the data.  
For example, a user queries a row and then later queries the same row, only to discover that the data has changed.

# Types of Read Operations

## 3. Phantom Reads

- A transaction re-runs a query returning a set of rows that satisfies a search condition and finds that another committed transaction has inserted additional rows that satisfy the condition.
- For example, a transaction queries the number of employees. Five minutes later it performs the same query, but now the number has increased by one because another user inserted a record for a new hire. More data satisfies the query criteria than before, but unlike in a fuzzy read the previously read data is unchanged.

# Isolation Levels

- **Read Uncommitted** will read uncommitted data from other transactions.
  - The database **does not place any locks on the data**, which increases transaction performance but at the cost of data consistency.
- **Read Committed** forces transactions to read only committed data.
  - **Default mode** of operation for most databases (incl. Oracle and SQL Server).
  - The database will **use exclusive locks on data**, causing other transactions to **wait** until the original transaction commits.

# Isolation Levels

- The **Repeatable Read** isolation level ensures that queries return consistent results.
  - Uses **shared locks** - ensure other transactions do not update a row after the original query reads it.
  - However, new rows are read (**phantom read**) as these rows did not exist when the first query ran.
- The **Serializable** isolation level is the most restrictive level defined by the ANSI SQL standard.
  - Deadlocks are always possible.
  - Most databases use a deadlock detection approach to transaction management, and, therefore, they will detect “deadlocks” during the transaction validation phase and reschedule the transaction.

# Preventable Read Phenomena by Isolation Level

**Table 10.15** Transaction Isolation Levels

	Isolation Level	Allowed			Comment
		Dirty Read	Nonrepeatable Read	Phantom Read	
Less restrictive  	Read Uncommitted	Y	Y	Y	Reads uncommitted data, and allows nonrepeatable reads and phantom reads.
	Read Committed	N	Y	Y	Does not allow uncommitted data reads but allows nonrepeatable reads and phantom reads.
	Repeatable Read	N	N	Y	Only allows phantom reads.
	Serializable	N	N	N	Does not allow dirty reads, nonrepeatable reads, or phantom reads.
Oracle/SQL Server Only	Read Only/Snapshot	N	N	N	Supported by Oracle and SQL Server. The transaction can only see the changes that were committed at the time the transaction started.

# Purpose of Isolation Levels

- The reason for the different levels of isolation is to increase transaction concurrency.
- The isolation levels go from the least restrictive (*Read Uncommitted*) to the more restrictive (*Serializable*).
- The **higher the isolation level the more locks (shared and exclusive) are required** to improve data consistency, at the expense of transaction concurrency performance.

# Setting Isolation Levels in DBMS

- Set at transaction level:

```
SET TRANSACTION ISOLATION LEVEL SERIALIZABLE;  
SET TRANSACTION ISOLATION LEVEL READ COMMITTED;
```

- Set at session level:

```
ALTER SESSION SET ISOLATION_LEVEL = SERIALIZABLE;  
ALTER SESSION SET ISOLATION_LEVEL = READ COMMITTED;
```

- MySQL and SQL Server support all four ANSI isolation levels.
- Oracle supports READ COMMITTED and SERIALIZABLE (Oracle would suggest that it has superior implementation of the other two isolation levels. By default, set to read committed)

# Locking Protocol

- **Locking Protocol**
  - Other approaches : *serializability, time-stamping, and shadow-paging.*
- If risk of interference is low => **Two-Phase Locking** ~ common approach although it requires *deadlock avoidance!!*
- A transaction can lock a database item from being updated (by another transaction) while it is being used. It acquires a lock on the item of interest and releases it when it has finished.
- Lock applies to a tuple :
  - **exclusive (write; X)**
  - **shared(read; S)**

# Exclusive VS Shared Locks

## Exclusive Lock

- An **exclusive lock** exists when access is reserved specifically for the transaction that locked the object.
- The exclusive lock must be used when the potential for conflict exists.
- An exclusive lock is issued when a transaction wants **to update (write)** a data item and *no locks are currently held on that data item by any other transaction.*

## Shared Lock

- A **shared lock** exists when concurrent transactions are granted read access on the basis of a common lock.
- No conflict as long as all the concurrent transactions are read-only.
- A shared lock is issued when a transaction wants to **read data** from the database and *no exclusive lock is held on that data item.*

# Problem: Lost Update

Time	User 1 (Trans A)	User2 (Trans B)
1	Retrieve t	
2		Retrieve t
3	Update t	
4		Update t
5		
6		
7		

t : tuple in a table.  
Trans A loses an update at t4.  
The update at t3 is lost (overwritten) at t4 by B.

# Solved: Lost Update

Time	User 1 (Trans A)	User2 (Trans B)
1	Retrieve t (get S-lock on t)	
2		Retrieve t (get S-lock on t)
3	Update t (request X-lock on t)	
4	wait	Update t (request X-lock on t)
5	wait	wait
6	wait	wait
7		

- No update lost but → Deadlock

# Problem: Uncommitted Dependency

- 2 Problems (T1-3 ; T6-8).
- One trans is allowed to retrieve/update) a tuple updated by another, but not yet committed.
- Trans A is dependent at time t2 on an **uncommitted change** made by Trans B, which is lost on Rollback.

Time	User 1 (Trans A)	User 2 (Trans B)
1		Update t
2	Retrieve t	
3		Rollback
4		
5		
6		Update t
7	Update t	
8		Rollback

# Solved: Uncommitted Dependency

Time	User 1 (Trans A)	User 2 (Trans B)
1		Update t (get X-lock on t)
2	Retrieve t (request S-lock on t)	-
3	wait	-
4	wait	-
5	wait	Commit / Rollback (releases X-lock on t)
6	Resume: Retrieve t (get S-lock on t)	
7	-	
8		

# Problem: Inconsistent Analysis

- Trans A sees inconsistent DB state after B updated Accumulator => performs inconsistent analysis.

Initially: Acc 1 = 40;      Acc2 = 50;      Acc3 = 30;

Time	User 1 (Trans A)	User 2 (Trans B)
1	Retrieve Acc 1 : Sum = 40	
2	Retrieve Acc2 : Sum = 90	
3		Retrieve Acc3 :
4		Update Acc3: 30 → 20
5		Retrieve Acc1:
6		Update Acc1: 40 → 50
7		commit
8	Retrieve Acc3: Sum = 110 (not 120)	

# Solved: Inconsistent Analysis

Time	User 1 (Trans A)	User 2 (Trans B)
1	Retrieve Acc1 : (get S-lock) Sum = 40	
2	Retrieve Acc2 : (get S-lock) Sum = 90	
3		Retrieve Acc3: (get S-lock)
4		Update Acc3: (get X-lock) 30 → 20
5		Retrieve Acc1: (get S-lock)
6		Update Acc1: (request X-lock) wait
7	Retrieve Acc3: (request S-lock) wait	wait wait wait

- Inconsistent Analysis is prevented → Deadlock