

MongoDB Queries I

CS 341 Database Systems

MongoDB Query API

- The way you will interact with your data.
- Used in 2 ways
 - CRUD Operations
 - Aggregation Pipelines

MongoDB Query API Uses

You can use the MongoDB Query API to perform:

- Adhoc queries with mongosh, Compass, VS Code, or a MongoDB driver for the programming language you use.
- Data transformations using aggregation pipelines.
- Document join support to combine data from different collections.
- Graph and geospatial queries.
- Full-text search.
- Indexing to improve MongoDB query performance.
- Time series analysis.

Install MongoDB

Create a
MongoDB
Atlas Account

MongoDB Mongosh (Shell)

- **db**
 - Shows the database you are using
- **show dbs**
 - Shows all available databases
- **use database_name**
 - Change or create a database
 - Example: use iba_db
 - Remember: In MongoDB, a database is not actually created until it gets content!

```
Atlas atlas-9egitm-shard-0 [primary] test> use iba_db
switched to db iba_db
Atlas atlas-9egitm-shard-0 [primary] iba_db> show dbs
mydb      48.00 KiB
test       72.00 KiB
admin     288.00 KiB
local     42.56 GiB
```

- Once a collection is created, the database gets some content and is now created.

```
Atlas atlas-9egitm-shard-0 [primary] iba_db> db.createCollection("restaurants")
{ ok: 1 }
Atlas atlas-9egitm-shard-0 [primary] iba_db> show dbs
iba_db      8.00 KiB
mydb      48.00 KiB
test       72.00 KiB
admin     288.00 KiB
local     42.56 GiB
Atlas atlas-9egitm-shard-0 [primary] iba_db> -
```

Create a Collection

- **db.createCollection("nameofCollection")**
 - Creates a new collection
- **db.students.insertOne(object)**
 - Collection during the insert process

Insert Documents

insertOne()

```
db.posts.insertOne({ title:  
  "Post Title 1", body: "Body  
  of post.", category:  
  "News", likes: 1, tags:  
  ["news", "events"], date:  
  Date() })
```

insertMany()

```
db.posts.insertMany([ {  
  title: "Post Title 2",  
  body: "Body of post.",  
  category: "Event", likes:  
  2, tags: ["news",  
  "events"], date: Date() },  
  { title: "Post Title 3",  
  body: "Body of post.",  
  category: "Technology",  
  likes: 3, tags: ["news",  
  "events"], date: Date() }  
])
```

Data Retrieval

- **db.collection.find()**
 - Retrieves data
 - Fetches all documents if left empty
- **db.collection.findOne()**

db.collection.find()

- Enclose in brackets
- First {} contains the query criteria
- Second {} specifies the projected columns (either 1s or 0s, only exception is _id which acts as the PK for the object).
- Within the query criteria, the comparator operators are to be used with :

SQL VS Mongo Query API

SQL	Mongo Query
Select * from Students	db.students.find()
Select student_name from Students	db.students.find({}, {"student_name":1})
Select student_name from Students Where major = 'CS'	db.students.find({"major": "CS"}, {"student_name":1})

Projection

```
Select student_name  
from Students  
Where major = 'CS'
```

```
db.students.find(  
  {"major": "CS"}, {"student_name": 1,  
   "_id":0})
```

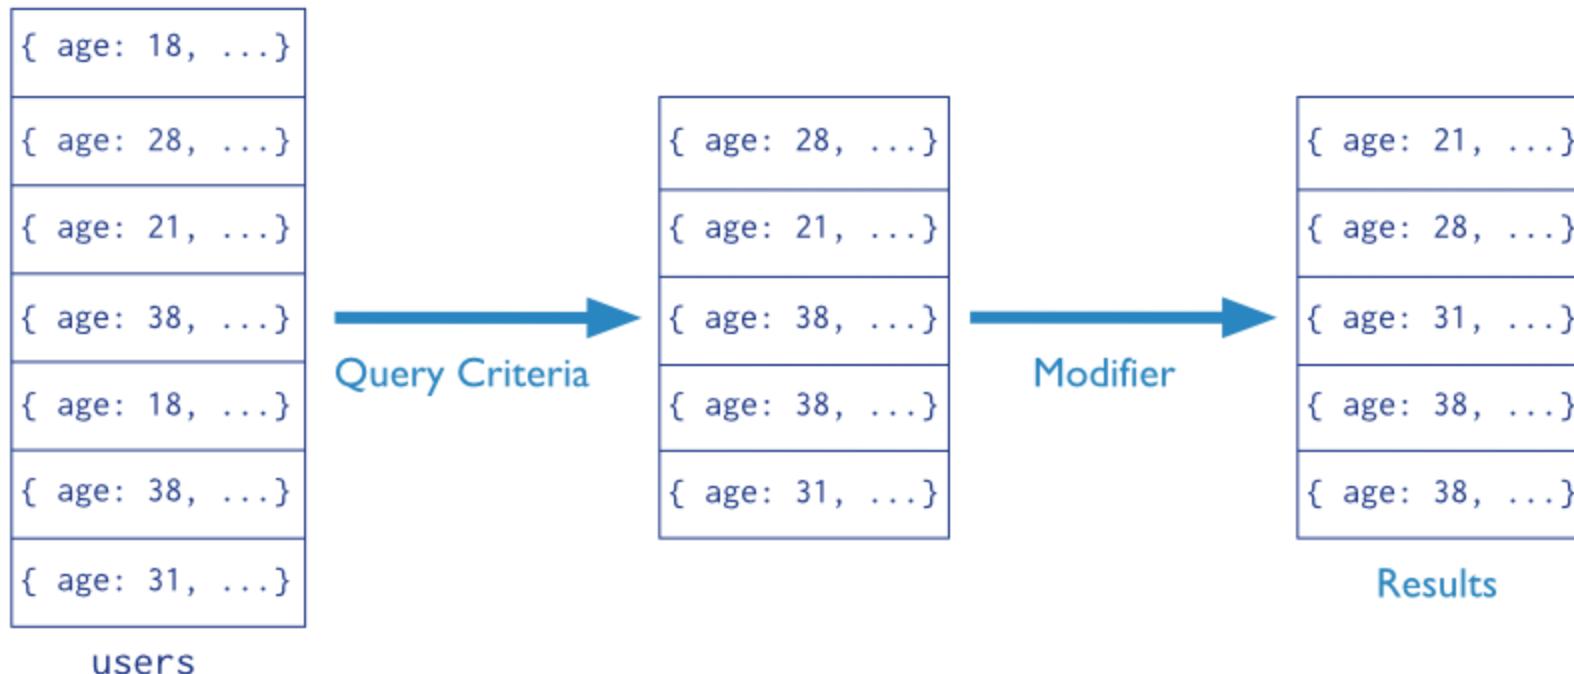
- Note: The `_id` field is also included. This field is always included unless specifically excluded.
- We use a 1 to include a field and 0 to exclude a field.
- You cannot use both 0 and 1 together except in case of `"_id"`
- Either specify the fields to include or to exclude.

Comparison Operators

- **\$eq** Values are equal
- **\$ne** Values are not equal
- **\$gt** Value is greater than another value
- **\$gte** Value is greater than or equal to another value
- **\$lt** Value is less than another value
- **\$lte** Value is less than or equal to another value
- **\$in** Value is matched within an array

Find the users of age greater than 18 and sort by age.

Collection Query Criteria Modifier
`db.users.find({ age: { $gt: 18 } }).sort({age: 1})`



Logical Operators

- **\$and** Returns documents where both queries match
- **\$or** Returns documents where either query matches
- **\$nor** Returns documents where both queries fail to match
- **\$not** Returns documents where the query does not match

Note: In MongoDB, when you provide multiple conditions in the find method, it implicitly interprets them as an AND operation. So, the use of \$and is often optional but recommended.

Logical Operators

- db.students.find({gender: "m", nationality: "Pakistan"})

Is equivalent to:

- db.students.find({\$and: [{ gender: "m" }, { nationality: "Pakistan" }]})
- db.musicians.find({\$or: [{ instrument: "Drums" }, { born: 1945 }] })
- db.musicians.find({ "instrument": { \$in: ["Keyboards", "Bass"] } })

Evaluation Operators

- **\$regex** Allows the use of regular expressions when evaluating field values
- **\$text** Performs a text search
- **\$where** Uses a JavaScript expression to match documents

Regex - contains

- All products that have t in them

```
db.purchase_orders.find( { product: { $regex: "t"} } )
```

- All products that have t in them (case insensitive)

```
db.purchase_orders.find( { product: { $regex: /t/i} } )
```

Regex - starts with

- Product starting with t:

```
db.purchase_orders.find( { product: { $regex: "^t" } } )
```

- Product starting with T (case insensitive)

```
db.purchase_orders.find( { product: { $regex: /^t/i } } )
```

Regex - ends with

Product ending with h:

```
db.purchase_orders.find( { product: { $regex: "h$"} } )
```

Case Insensitive:

```
db.purchase_orders.find( { product: { $regex: /h$/i} } )
```

Regex - start and end

```
db.purchase_orders.find({product:{$regex: /^M.*e$/}})
```

`.*` → One or Many

- `.` : Matches any single character (except for a newline character).
- `*` : Matches zero or more occurrences of the preceding character or group.

Sort

- Ascending **1**
- Descending **-1**

```
db.students.find({gender: 'm'}).sort({nationality: -1});
```

```
db.students.find({gender: 'm'}).sort({nationality: -1,  
firstName: 1});
```

Limit

- Limit records to just 2

```
db.students.find({gender: 'f', $or: [{nationality: 'pakistani'}, {nationality: 'american'}]}).limit(2);
```

- Retrieve the 3rd and 4th record by skipping over first 2.

```
db.students.find({gender: 'f', $or: [{nationality: 'pakistani'}, {nationality: 'american'}]}).limit(2).skip(2);
```

Count

Counting number of batteries sold

```
db.purchase_orders.countDocuments({product: "battery"})
```

Count number of documents found

```
db.purchase_orders.find({}).count()
```

Distinct

- List all the distinct products

```
db.purchase_orders.distinct("product")
```

ElemMatch

Element
Matching

ElemMatch for Querying Arrays

- The \$elemMatch operator matches documents that contain an array field with **least one element that matches all the specified query criteria.**
- This is primarily used when you want to *match documents based on conditions within an array field.*
- It is useful when you want **all the specified conditions** to be met by **a single element within an array.**
- You don't need to use the \$elemMatch operator when specifying a single search condition on an array field. You may directly access the field through dot notation.

ElemMatch in Scores collection

```
{ _id: 1, results: [ 82, 85, 88 ] }  
{ _id: 2, results: [ 75, 88, 89 ] }
```

The following query matches only those documents where the results array contains at least one element that is **both** greater than or equal to 80 and is less than 85:

```
db.scores.find(  
  { results: { $elemMatch: { $gte: 80, $lt: 85 } } }  
)
```

Dot Notation

- Dot notation is used to access fields within nested documents or arrays.
- It helps you navigate through the structure of a document.
- Example *Results.subject.score*

Restaurants Collection

```
{  
  "address": {  
    "building": "1007",  
    "coord": [ -73.856077, 40.848447 ],  
    "street": "Morris Park Ave",  
    "zipcode": "10462"  
  },  
  "borough": "Bronx",  
  "cuisine": "Bakery",  
  "grades": [  
    { "date": { "$date": 1393804800000 }, "grade": "A", "score": 2 },  
    { "date": { "$date": 1378857600000 }, "grade": "A", "score": 6 },  
    { "date": { "$date": 1358985600000 }, "grade": "A", "score": 10 },  
    { "date": { "$date": 1322006400000 }, "grade": "A", "score": 9 },  
    { "date": { "$date": 1299715200000 }, "grade": "B", "score": 14 }  
  ],  
  "name": "Morris Park Bake Shop",  
  "restaurant_id": "30075445"  
}
```

Display restaurant names with score more than 80 but less than 100.

date: ISODate("2014-06-27T00:00:00.000Z")

```
db.restaurants.find({ "grades.score": { $gt: 80, $lt :100}});
```

```
[  
  {  
    _id: ObjectId("65591b9d493a114effe17225"),  
    address: {  
      building: '65',  
      coord: [ -73.9782725, 40.7624022 ],  
      street: 'West 54 Street',  
      zipcode: '10019'  
    },  
    borough: 'Manhattan',  
    cuisine: 'American',  
    grades: [  
      {  
        date: ISODate("2014-08-22T00:00:00.000Z"),  
        grade: 'A',  
        score: 11  
      },  
      {  
        date: ISODate("2014-03-28T00:00:00.000Z"),  
        grade: 'C',  
        score: 131  
      },  
      {  
        date: ISODate("2013-09-25T00:00:00.000Z"),  
        grade: 'A',  
        score: 11  
      },  
      {  
        date: ISODate("2013-09-25T00:00:00.000Z"),  
        grade: 'A',  
        score: 11  
      }  
    ]  
  }]  
]
```

```
        },
        {
            date: ISODate("2013-04-08T00:00:00.000Z"),
            grade: 'B',
            score: 25
        },
        {
            date: ISODate("2012-10-15T00:00:00.000Z"),
            grade: 'A',
            score: 11
        },
        {
            date: ISODate("2011-10-19T00:00:00.000Z"),
            grade: 'A',
            score: 13
        }
    ],
    name: "Murals On 54/Randolph's S",
    restaurant_id: '40372466'
},
```

```
Atlas atlas-9egitm-shard-0 [primary] iba_db> db.restaurants.find({ "grades": { $elemMatch: { "score": { $gt: 80, $lt: 100 } } } }, { "name": 1 }).count()
3
Atlas atlas-9egitm-shard-0 [primary] iba_db> db.restaurants.find({ "grades.score": { $gt: 80, $lt: 100 } }, { "name": 1 }).count()
4
Atlas atlas-9egitm-shard-0 [primary] iba_db>
```



- **db.restaurants.find({grades : {
\$elemMatch:{"score":{\$gt : 80 ,
\$lt :100}}}});**

```
{
  _id: ObjectId("673ceb13d4e1d7875c6d112e"),
  address: {
    building: '',
    coord: [ -74.0163793, 40.7167671 ],
    street: 'Hudson River',
    zipcode: '10282'
  },
  borough: 'Manhattan',
  cuisine: 'American',
  grades: [
    {
      date: ISODate("2014-06-27T00:00:00.000Z"),
      grade: 'C',
      score: 89
    },
    {
      date: ISODate("2013-06-06T00:00:00.000Z"),
      grade: 'A',
      score: 6
    },
    {
      date: ISODate("2012-06-19T00:00:00.000Z"),
      grade: 'A',
      score: 13
    }
  ],
  name: 'West 79Th Street Boat Basin Cafe',
  restaurant_id: '40756344'
}
```

elemMatch is more Specific

- **Directly on "grades.score":**

```
db.restaurants.find({ "grades.score": { $gt: 80, $lt: 100 } }, { "name": 1 })
```

- This query will return documents where there is at least one grade within the "grades" array with a score in range 80 and less than 100. However, it doesn't ensure that both conditions are satisfied by the same grade element within the array.

- **Using \$elemMatch:**

```
db.restaurants.find({ "grades": { $elemMatch: { "score": { $gt: 80, $lt: 100 } } } }, { "name": 1 })
```

- This query will return documents where there is at least one grade within the "grades" array that has a score greater than 80 and less than 100. It ensures that a single grade within the array satisfies both conditions simultaneously.

\$all operator

- \$all is used to match array fields where the array contains all specified elements. It can also be combined with \$elemMatch to apply specific conditions to elements in the array.
- *Q. Find documents where there is at least one grade with score 10 and at least one grade with grade 'A'*
- Notice in our dataset, a score of 10 would imply a grade A hence the condition would appear to be correct in a single array element.
Let's consider a smaller and different version of the restaurants collection

```
{  
    "name": "Restaurant A",  
    "grades": [  
        { "score": 10, "grade": "B" },  
        { "score": 8, "grade": "A" },  
        { "score": 5, "grade": "C" }  
    ]  
},  
{  
    "name": "Restaurant B",  
    "grades": [  
        { "score": 10, "grade": "A" },  
        { "score": 6, "grade": "B" }  
    ]  
},  
{  
    "name": "Restaurant C",  
    "grades": [  
        { "score": 7, "grade": "B" },  
        { "score": 8, "grade": "C" }  
    ]  
}
```

```
db.restaurants.find({  
    "grades": {  
        $all: [  
            { $elemMatch: { "score": 10 } },  
            { $elemMatch: { "grade": "A" } }  
        ]  
    }  
})
```

- When using `$all` with `$elemMatch`, we are specifying that the array must contain elements such that both conditions are satisfied, but the conditions do not need to be met by the same element.

```
[  
    { "name": "Restaurant A" },  
    { "name": "Restaurant B" }  
]
```